

AI-Powered Enterprise Knowledge Base (RAG)

Document Ingestion, AWS OpenSearch Vector Search & AWS Bedrock LLMs

- Python 3.11
- FastAPI
- LangChain
- AWS Bedrock
- AWS OpenSearch Serverless
- Amazon S3

1. PROJECT OVERVIEW & ARCHITECTURE OBJECTIVES


Enterprise knowledge management requires real-time query resolution across private PDFs and corporate policy documents while keeping data secured within an AWS cloud boundary. This project details the design and deployment of an **AI-Powered Enterprise Knowledge Base (RAG System)**.

When documents are uploaded to an **Amazon S3 Bucket**, S3 event notifications automatically trigger a serverless **AWS Lambda** worker. The worker extracts text using **LangChain**, creates vector embeddings via **AWS Bedrock Titan Embeddings**, and indexes them into **AWS OpenSearch Serverless**. A **FastAPI** web backend retrieves top context chunks and synthesizes natural language answers using **AWS Bedrock Claude 3 / Titan LLMs**.

Architecture Layer	AWS Service / Tech	Technical Function & Purpose
Document Ingestion	Amazon S3 Bucket	Stores PDF/Doc assets; triggers automated event-driven processing.
Chunking & Vectorization	AWS Lambda & Bedrock Titan	LangChain text splitter generates 1536-dim embeddings via Titan.
Vector Database	AWS OpenSearch Serverless	k-NN vector search engine storing vector index for semantic retrieval.
Query API & Retrieval	FastAPI & LangChain	Asynchronous REST API orchestrating context search and prompt building.
LLM Synthesis	AWS Bedrock (Claude 3 / Titan)	Generates ground-truth natural language answers without hallucinations.

2. SYSTEM ARCHITECTURE & COMPONENT WORKFLOW DIAGRAM

FIGURE 1: AWS BEDROCK & OPENSEARCH ENTERPRISE RAG ARCHITECTURE



3. CORE IMPLEMENTATION SNIPPET (FASTAPI & BEDROCK)

```
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
import boto3
import json

app = FastAPI(title="AWS Bedrock & OpenSearch Enterprise RAG")
bedrock_runtime = boto3.client('bedrock-runtime', region_name='us-east-1')

class QueryRequest(BaseModel):
    question: str

@app.post("/query")
def query_knowledge_base(payload: QueryRequest):
    # 1. Retrieve Context from OpenSearch (Simulated context)
    context = "AWS OpenSearch Serverless supports vector search with k-NN indices."

    # 2. Build Prompt for AWS Bedrock
    prompt = f"Human: Context:\n{context}\n\nQuestion: {payload.question}\n\nAssistant:"

    body = json.dumps({
        "prompt": prompt,
        "max_tokens_to_sample": 500,
        "temperature": 0.2
    })

    response = bedrock_runtime.invoke_model(
        modelId="anthropic.claude-v2",
        body=body
    )
    result = json.loads(response['body'].read().decode('utf-8'))
    return {"question": payload.question, "answer": result.get("completion")}
```

4. KEY TECHNICAL ACHIEVEMENTS

Enterprise Security & Performance Impact

- **Zero Data Exposure:** All document vectors and LLM invocations remain enclosed inside AWS security boundaries.
- **Sub-Second k-NN Search:** AWS OpenSearch Serverless delivers sub-100ms vector similarity lookup.
- **Fully Serverless Auto-Scaling:** Ingests thousands of pages automatically without managing server clusters.